



• WUST •

学习总结



Mode by : 董勇

目录

CONTENT

01

论文



02

视频



03

实验



PART

01

S. Han et al., “ESE: Efficient Speech Recognition Engine with Sparse LSTM on FPGA,” in Proceedings of the 2017 ACM/SIGDA International Symposium on Field-Programmable Gate Arrays, Monterey California USA, Feb. 2017, pp. 75–84, doi: 10.1145/3020078.3021745.

01 Paper

motivation


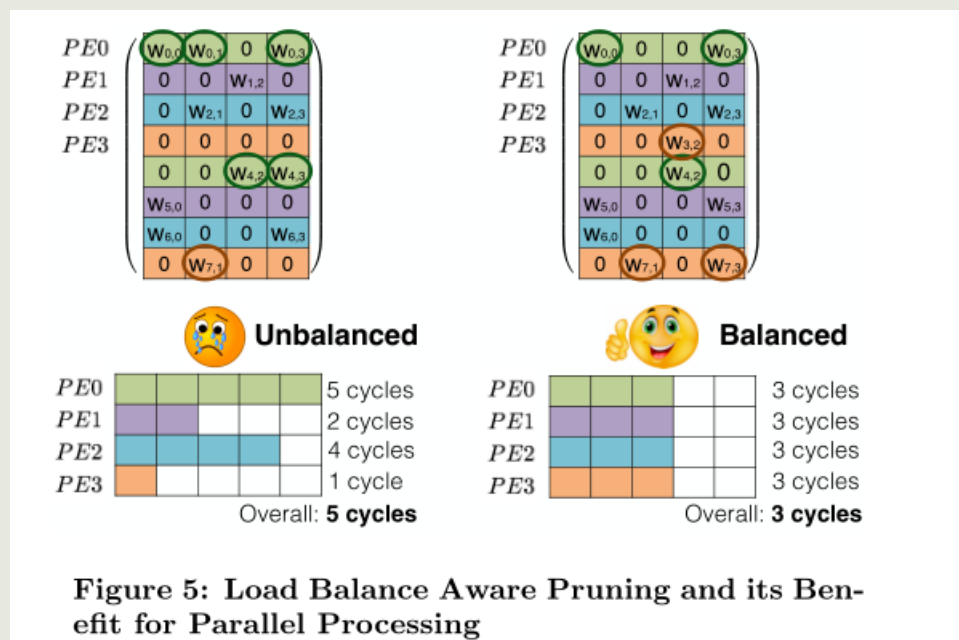
LSTM模型过大：
存储消耗
运算消耗

contribution

- 剪枝与量化应用于LSTM，剪枝时用到了负载均衡 (load balance)，将LSTM带来了大量的压缩20x (剪枝10x，量化2x)
- 提出了一个scheduler，用于编码和分组相应的复杂的压缩的LSTM，分配给PE

01 Paper

Pruning & Quantization



将分组了的参数按照一致的比例去稀疏，而不是原来那样全局稀疏；
并通过retraining把损失的精度补回来。这样就做到了负载均衡的稀疏参数了。

量化：压缩32bit的浮点数到12bit定点，然后运用线性的量化来实现于weight和activation

01 Paper

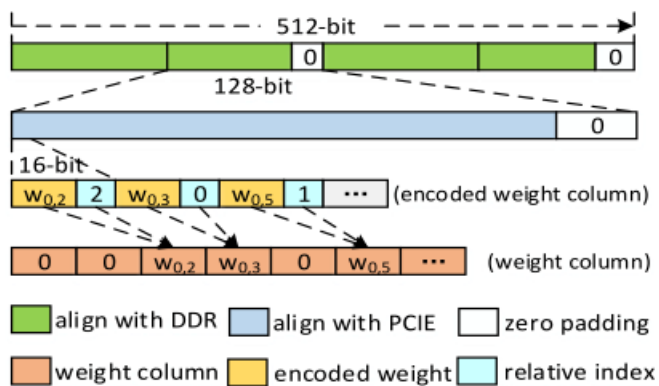


Figure 7: Encoding in CSC format and data align using zero-padding.

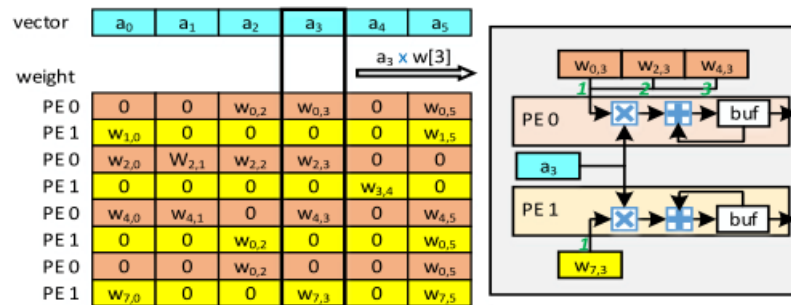


Figure 8: The computation pattern: non-zero weights in a column are assigned to 2 PEs, and every PE multiply-add their weights with the same element from the shared vector.

01 Paper

Architecture Overview

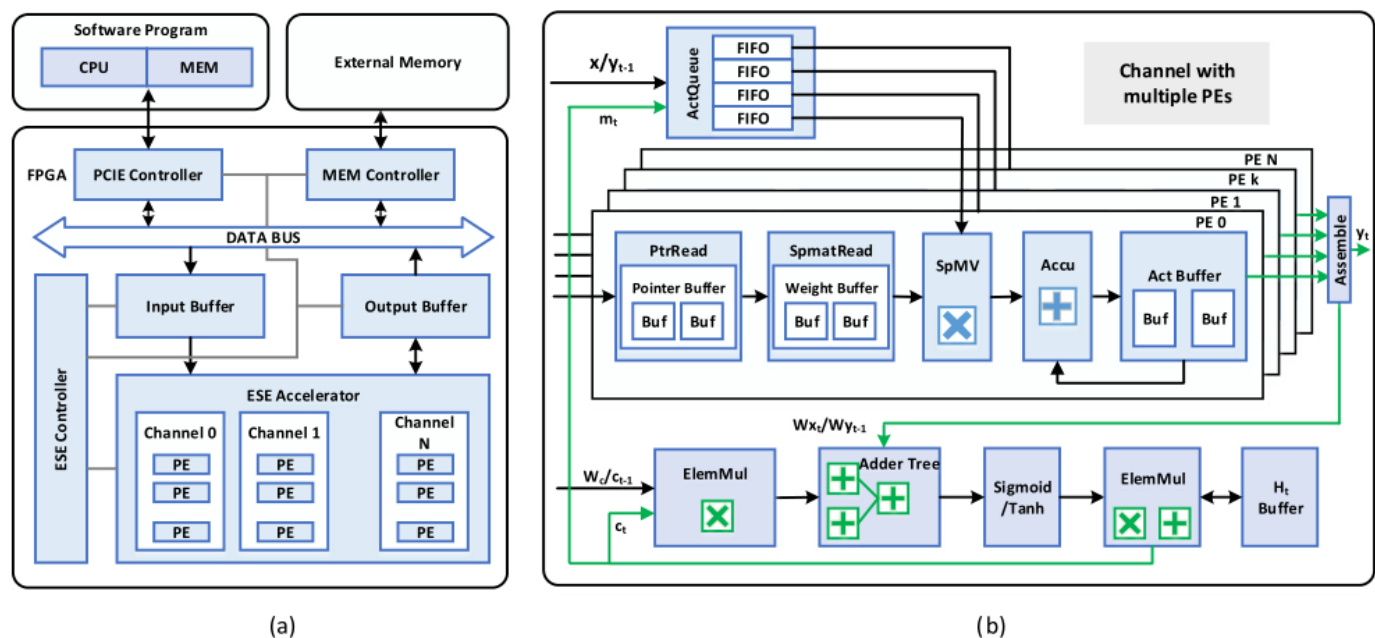


Figure 10: The Efficient Speech Recognition Engine (ESE) accelerator system: (a) the overall ESE system architecture; (b) one channel with multiple processing elements (PEs).

ESE Channel Architecture:

- **ActQueue**: 每个PE之前设置一个队列, 用于存储, 这样PE之间不同同步, 只用处理各自队列上的值。
- **SpmatRead**: 稀疏矩阵读取单元, 用于编码权重矩阵的存储和输出。
- **SpMV**: 矩阵相乘单元。
- **ElemMul**: 元素级别的乘法单元。
- **Adder Tree**: 累加结果和偏置bias

01 Paper

Table 7: Power consumption of different platforms.

Platform	CPU Dense	CPU Sparse	GPU Dense	GPU Sparse	ESE
Power	111W	38W	202W	136W	41W

Table 8: Performance comparison of running LSTM on ESE, CPU and GPU

Plat.	ESE on FPGA (ours)									CPU		GPU	
Matrix	Matrix Size	Sparsity (%) ¹	Compres. Matrix (Bytes) ²	Theoreti. Comput. Time (μs)	Real Comput. Time (μs)	Total Operat. (GOP)	Real Perform. (GOP/s)	Equ. Operat. (GOP)	Equ. Perform. (GOP/s)	Real Comput. Time (μs)		Real Comput. Time (μs)	
										Dense	Sparse	Dense	Sparse
W_{ix}	1024×153	11.7	36608	2.9	5.36	0.0012	218.6	0.010	1870.7	1518.4 ³	670.4	34.2	58.0
W_{fx}	1024×153	11.7	36544	2.9	5.36	0.0012	218.2	0.010	1870.7				
W_{cx}	1024×153	11.8	37120	2.9	5.36	0.0012	221.6	0.010	1870.7				
W_{ox}	1024×153	11.5	35968	2.8	5.36	0.0012	214.7	0.010	1870.7				
W_{ir}	1024×512	11.3	118720	9.3	10.31	0.0038	368.5	0.034	3254.6	3225.0 ⁴	2288.0	81.3	166.0
W_{fr}	1024×512	11.5	120832	9.4	10.01	0.0039	386.3	0.034	3352.1				
W_{cr}	1024×512	11.2	117760	9.2	9.89	0.0038	381.2	0.034	3394.5				
W_{or}	1024×512	11.5	120256	9.4	10.04	0.0038	383.5	0.034	3343.7				
W_{ym}	512×1024	10.0	104832	8.2	15.66	0.0034	214.2	0.034	2142.7	1273.9	611.5	124.8	63.4
Total	3248128	11.2	728640	57.0	82.7	0.0233	282.2	0.208	2515.7	6017.3	3569.9	240.3	287.4

CPU sparse比dense快，因为CPU擅长串行处理，而GPU sparse比dense慢，因为GPU注重吞吐量。对于稀疏的LSTM, CPU和GPU都更快，因为内存带宽的节省更加显著。

Conclusion:

ESE在sparse LSTM取得了282 GOPS处理速率，相当于非稀疏的2.52 TOPS的处理速率。

在语音识别的数据集上ESE取得了比 i7 5930和titan X GPU快43x和3x的速度提升

比CPU和GPU能耗降低40x和11.5x

PART

02

- transformer
- NLP(1-2)

02/ Video

1. CNN在文本中的编码能力弱于RNN，而RNN是序列模型，并行能力差，计算缓慢，而且只能考虑一个方向上的信息。Transformer可以综合的考虑两个方向的信息，而且有非常好的并行性质
2. Transformer也包含大量的矩阵运算，cnn中矩阵运算优化的方法也同样适用

PART

03

- RNN
-

	A	B	C	
	id	label		
	0	0		
	1	0		
	2	0		
	3	1		
	4	1		
	5	1		
	6	0		
	7	1		
0	8	1		
1	9	1		
2	10	1		
3	11	1		
4	12	1		
5	13	0		
6	14	1		
7	15	1		
8	16	0		
9	17	0		
0	18	1		
1	19	0		
2	20	1		
3	21	1		
4	22	0		
5	23	0		
6	24	1		
7	25	0		
8	26	0		
9	27	1		
0	28	0		

03/ Experiment

文本情感分析，训练文件中的句子被标记为“1”或“0”，分别对应句子的情感色彩是“负面”与“正面”

```
import os
import numpy as np
import pandas as pd
import argparse
from gensim.models import word2vec

def train_word2vec(x):
    # 訓練 word to vector 的 word embedding
    model = word2vec.Word2Vec(x, size=250, window=5, min_count=5, workers=12, iter=10, sg=1)
    return model

if __name__ == "__main__":
    print("loading training data ...")
    train_x, y = load_training_data('training_label.txt')
    train_x_no_label = load_training_data('training_nolabel.txt')

    print("loading testing data ...")
    test_x = load_testing_data('testing_data.txt')

    #model = train_word2vec(train_x + train_x_no_label + test_x)
    model = train_word2vec(train_x + test_x)

    print("saving model ...")
    # model.save(os.path.join(path_prefix, 'model/w2v_all.model'))
    model.save(os.path.join(path_prefix, 'w2v_all.model'))
```

```
from torch import nn
from gensim.models import Word2Vec

class Preprocess():
    def __init__(self, sentences, sen_len, w2v_path="./w2v.model"):
        self.w2v_path = w2v_path
        self.sentences = sentences
        self.sen_len = sen_len
        self.idx2word = []
        self.word2idx = {}
        self.embedding_matrix = []

    def get_w2v_model(self):
        # 把之前訓練好的 word to vec 模型讀進來
        self.embedding = Word2Vec.load(self.w2v_path)
        self.embedding_dim = self.embedding.vector_size

    def add_embedding(self, word):
        # 把 word 加進 embedding，並賦予他一個隨機生成的 representation vector
        # word 只會是 "<PAD>" 或 "<UNK>"
        vector = torch.empty(1, self.embedding_dim)
        torch.nn.init.uniform_(vector)
        self.word2idx[word] = len(self.word2idx)
        self.idx2word.append(word)
        self.embedding_matrix = torch.cat([self.embedding_matrix, vector], 0)

    def make_embedding(self, load=True):
        print("Get embedding ...")
        # 取得訓練好的 Word2vec word embedding
```

```
import torch
from torch import nn

class LSTM_Net(nn.Module):
    def __init__(self, embedding, embedding_dim, hidden_dim, num_layers, dropout=0.5, fix_embedding=True):
        super(LSTM_Net, self).__init__()
        # 製作 embedding layer
        self.embedding = torch.nn.Embedding(embedding.size(0), embedding.size(1))
        self.embedding.weight = torch.nn.Parameter(embedding)
        # 是否將 embedding fix 住，如果 fix_embedding 為 False，在訓練過程中，embedding 也會跟著被訓練
        self.embedding.weight.requires_grad = False if fix_embedding else True
        self.embedding_dim = embedding.size(1)
        self.hidden_dim = hidden_dim
        self.num_layers = num_layers
        self.dropout = dropout
        self.lstm = nn.LSTM(embedding_dim, hidden_dim, num_layers=num_layers, batch_first=True)
        self.classifier = nn.Sequential( nn.Dropout(dropout),
                                         nn.Linear(hidden_dim, 1),
                                         nn.Sigmoid() )

    def forward(self, inputs):
        inputs = self.embedding(inputs)
        x, _ = self.lstm(inputs, None)
        # x 的 dimension (batch, seq_len, hidden_size)
        # 取用 LSTM 最後一層的 hidden state
        x = x[:, -1, :]
        x = self.classifier(x)
        return x
```

