

20210210

孙阳

# 论文学习

- 《Efficient Speech Recognition Engine with Compressed LSTM on FPGA》 基于FPGA的高效压缩LSTM语音识别引擎 FPGA 2017
- 《Efficient methods and hardware for deep learning》
- 亮点：
  - 考虑到最终要多核运行并行加速的时候不同核心之间的负载均衡，提出了Load-balance-aware pruning算法
  - Deep compression，从软件端极大的压缩了网络的权重。
  - DSD，密集，稀疏，密集的训练方法，一种新的网络训练方法，能从训练层面一定的提升网络准确率
  - EIE: 在Deep compression 的基础上，EIE是基于硬件的稀疏网络加速实现，硬件上达到很好的效果。

# Deep compression

- 过程:

- 1. 剪枝
- 2. 稀疏矩阵存储
- 3. 权值共享

- 压缩率:

$$r = \frac{nb}{n \log_2(k) + kb}$$

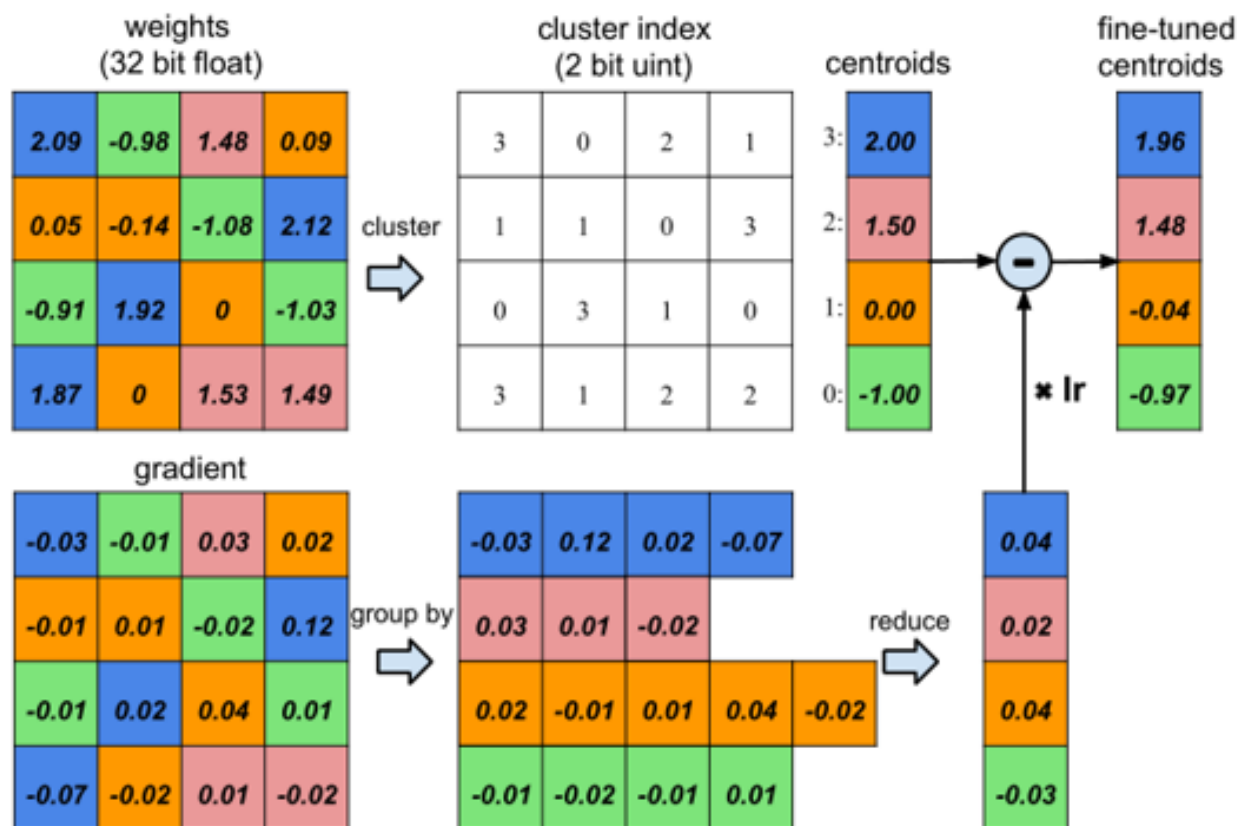


Figure 3: Weight sharing by scalar quantization (top) and centroids fine-tuning (bottom).

[https://blog.csdn.net/weixin\\_36474809](https://blog.csdn.net/weixin_36474809)

# DSD网络

- Dense-sparse-dense training for deep neural networks
- 先Dense训练，获得一个网络权重；然后spars将网络剪枝，然后继续dense训练。从而获得更好的性能。

---

**Algorithm 1:** Workflow of DSD training

---

**Initialization:**  $W^{(0)}$  with  $W^{(0)} \sim N(0, \Sigma)$

**Output :**  $W^{(t)}$ .

————— *Initial Dense Phase* —————

**while** not converged **do**

$W^{(t)} = W^{(t-1)} - \eta^{(t)} \nabla f(W^{(t-1)}; x^{(t-1)});$

$t = t + 1;$

**end**

————— *Sparse Phase* —————

// initialize the mask by sorting and keeping the Top-k weights.

$S = \text{sort}(|W^{(t-1)}|); \lambda = S_{k_i}; \text{Mask} = \mathbb{1}(|W^{(t-1)}| > \lambda);$

**while** not converged **do**

$W^{(t)} = W^{(t-1)} - \eta^{(t)} \nabla f(W^{(t-1)}; x^{(t-1)});$

$W^{(t)} = W^{(t)} \cdot \text{Mask};$

$t = t + 1;$

**end**

————— *Final Dense Phase* —————

**while** not converged **do**

$W^{(t)} = W^{(t-1)} - \eta^{(t)} \nabla f(W^{(t-1)}; x^{(t-1)});$

$t = t + 1;$

**end**

**goto** Sparse Phase for iterative DSD; [https://blog.csdn.net/weixin\\_36474809](https://blog.csdn.net/weixin_36474809)

---

# EIE网络

$$b_i = ReLU \left( \sum_{j=0}^{n-1} W_{ij} a_j \right)$$

$$b_i = ReLU \left( \sum_{j \in X_i \cap Y} S[I_{ij}] a_j \right)$$

# 理论学习

- 李宏毅NLP-BERT P17-P19

# 实验

- 一、利用cpu和HLS运行同一份RNN训练模型，进行运算速度对比
- 二、Deep-Compression-PyTorch
  - 实验使用了MNIST dataset 训练了LeNet-300-100 model
  - 重点：权值共享、哈夫曼编码

# 实验1代码

```
void train()
{
    int epoch, i, j, k, m, p;
    vector<double*> layer_1_vector;    //保存隐藏层
    vector<double> layer_2_delta;    //保存误差关于Layer 2 输出值的偏导

    for(epoch=0; epoch<10000; epoch++) //训练次数
    {
        double e = 0.0; //误差
        for(i=0; i<layer_1_vector.size(); i++)
            delete layer_1_vector[i];
        layer_1_vector.clear();
        layer_2_delta.clear();

        int d[binary_dim];    //保存每次生成的预测值
        memset(d, 0, sizeof(d));

        int a_int = (int)randval(largest_number/2.0); //随机生成一个加数 a
        int a[binary_dim];
        int2binary(a_int, a);    //转为二进制数

        int b_int = (int)randval(largest_number/2.0); //随机生成另一个加数 b
        int b[binary_dim];
        int2binary(b_int, b);    //转为二进制数

        int c_int = a_int + b_int;    //真实的和 c
        int c[binary_dim];
        int2binary(c_int, c);    //转为二进制数

        double *layer_1 = new double[hiddenode];
        for(i=0; i<hiddenode; i++)    //在0时刻是没有之前的隐含层的，所以初始化一个全为0的
            layer_1[i] = 0;
        layer_1_vector.push_back(layer_1);
    }
```

```
//正向传播
for(p=0; p<binary_dim; p++)    //循环遍历二进制数组，从最低位开始
{
    layer_0[0] = a[p];
    layer_0[1] = b[p];
    double y = (double)c[p];    //实际值
    layer_1 = new double[hiddenode]; //当前隐含层

    for(j=0; j<hiddenode; j++)
    {
        //输入层传播到隐含层
        double o1 = 0.0;
        for(m=0; m<innode; m++)
            o1 += layer_0[m] * w[m][j];

        //之前的隐含层传播到现在的隐含层
        double *layer_1_pre = layer_1_vector.back();
        for(m=0; m<hiddenode; m++)
            o1 += layer_1_pre[m] * wh[m][j];

        layer_1[j] = sigmoid(o1);    //隐藏层各单元输出
    }

    for(k=0; k<outnode; k++)
    {
        //隐藏层传播到输出层
        double o2 = 0.0;
        for(j=0; j<hiddenode; j++)
            o2 += layer_1[j] * w1[j][k];
        layer_2[k] = sigmoid(o2);    //输出层各单元输出
    }
}
```



# 实验2结果

```
--- Before pruning ---
fc1.weight      | nonzeros = 235200 / 235200 (100.00%) | total_pruned = 0 | shape = (300, 784)
fc1.bias        | nonzeros = 300 / 300 (100.00%) | total_pruned = 0 | shape = (300,)
fc2.weight      | nonzeros = 30000 / 30000 (100.00%) | total_pruned = 0 | shape = (100, 300)
fc2.bias        | nonzeros = 100 / 100 (100.00%) | total_pruned = 0 | shape = (100,)
fc3.weight      | nonzeros = 1000 / 1000 (100.00%) | total_pruned = 0 | shape = (10, 100)
fc3.bias        | nonzeros = 10 / 10 (100.00%) | total_pruned = 0 | shape = (10,)
alive: 266610, pruned : 0, total: 266610, Compression rate : 1.00x ( 0.00% pruned)
Pruning with threshold : 0.22420135140419006 for layer fc1
Pruning with threshold : 0.1908438801765442 for layer fc2
Pruning with threshold : 0.23130165040493011 for layer fc3
Test set: Average loss: 1.0954, Accuracy: 6761/10000 (67.61%)
--- After pruning ---
fc1.weight      | nonzeros = 10285 / 235200 ( 4.37%) | total_pruned = 224915 | shape = (300, 784)
fc1.bias        | nonzeros = 300 / 300 (100.00%) | total_pruned = 0 | shape = (300,)
fc2.weight      | nonzeros = 1360 / 30000 ( 4.53%) | total_pruned = 28640 | shape = (100, 300)
fc2.bias        | nonzeros = 100 / 100 (100.00%) | total_pruned = 0 | shape = (100,)
fc3.weight      | nonzeros = 69 / 1000 ( 6.90%) | total_pruned = 931 | shape = (10, 100)
fc3.bias        | nonzeros = 10 / 10 (100.00%) | total_pruned = 0 | shape = (10,)
alive: 12124, pruned : 254486, total: 266610, Compression rate : 21.99x ( 95.45% pruned)
--- Retraining ---
```

# 实验2结果

```
--- After Retraining ---
fc1.weight | nonzeros = 10285 / 235200 ( 4.37%) | total_pruned = 224915 | shape = (300, 784)
fc1.bias   | nonzeros = 300 / 300 (100.00%) | total_pruned = 0 | shape = (300,)
fc2.weight | nonzeros = 1360 / 30000 ( 4.53%) | total_pruned = 28640 | shape = (100, 300)
fc2.bias   | nonzeros = 100 / 100 (100.00%) | total_pruned = 0 | shape = (100,)
fc3.weight | nonzeros = 69 / 1000 ( 6.90%) | total_pruned = 931 | shape = (10, 100)
fc3.bias   | nonzeros = 10 / 10 (100.00%) | total_pruned = 0 | shape = (10,)
> alive: 12124, pruned : 254486, total: 266610, Compression rate : 21.99x ( 95.45% pruned)
```

Layer	original	compressed	improvement	percent
fc1.weight	83484	19452	4.29x	23.30%
fc1.bias	1200	1200	1.00x	100.00%
fc2.weight	11284	3196	3.53x	28.32%
fc2.bias	400	400	1.00x	100.00%
fc3.weight	596	398	1.50x	66.78%
fc3.bias	40	40	1.00x	100.00%
total	97004	24686	3.93x	25.45%